

Dimensinality Reduction

Main References

Ameet Talwalkar and Henry Chai, **ML with Large Datasets**, CMU

Outline

- Data Preprocessing
- Data Visualization
- Dimensionality Reduction
- PCA
- Distributed PCA

Data Pre-processing

- ETL (extract-transfer-load)
- Cleaning data
 - Missing features/labels
 - Duplicated observations
 - Formatting errors
- Understanding data
 - Summarization
 - Exploration
 - Visualization

What questions can you ask better understand the data?

- What are features and data types?
- What are the units?
- How many observations?
- How data is collected?
- Are features are correlated?
- ...

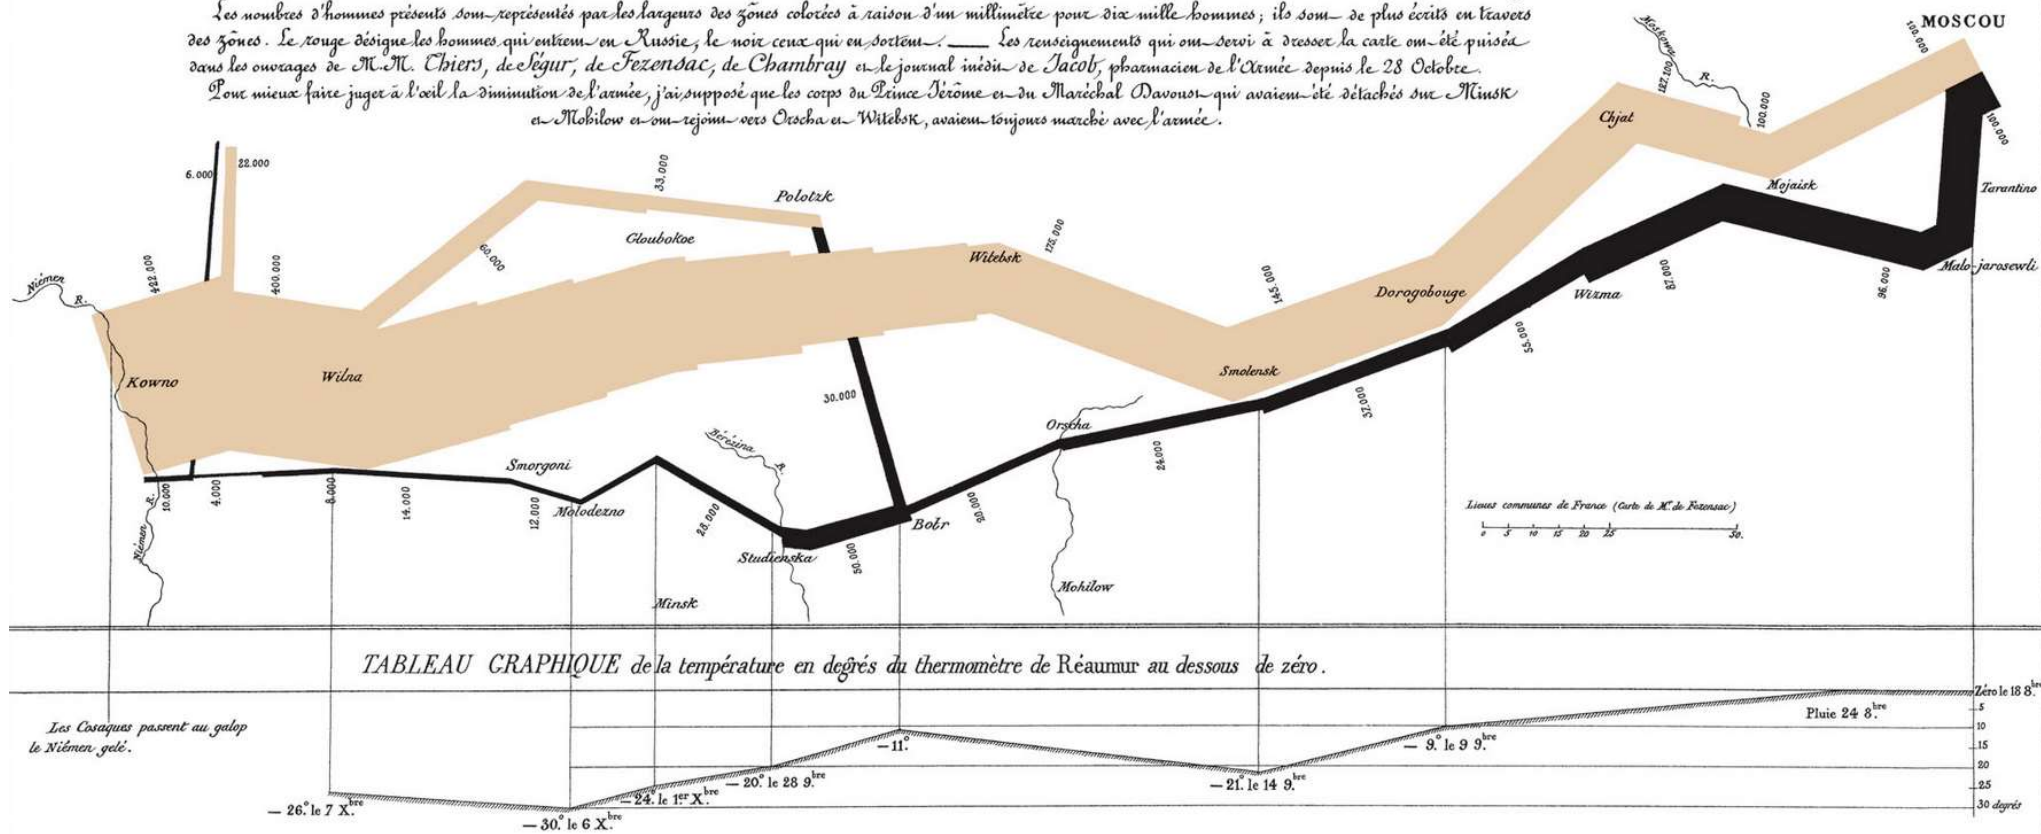
Data Visualization

- Visualizations can be used to
 - Provide insight about trends/groups/relationships
 - Reveal systematic errors
 - Aid in model selection
 - Evaluate training (e.g., measure convergence)
 - Interpret/explain predictions

Data Visualization

Carte Figurative des pertes successives en hommes de l'Armée Française dans la campagne de Russie 1812-1813.
 Dressée par M. Minard, Inspecteur Général des Ponts et Chaussées en retraite. Paris, le 20 Novembre 1869.

Les nombres d'hommes présents sont représentés par les largeurs des zones colorées à raison d'un millimètre pour dix mille hommes; ils sont de plus écrits en travers des zones. Le rouge désigne les hommes qui entrent en Russie, le noir ceux qui en sortent. Les renseignements qui ont servi à dresser la carte ont été puisés dans les ouvrages de M. M. Chiers, de Léger, de Fezensac, de Chambray et le journal inédit de Jacob, pharmacien de l'Armée depuis le 28 Octobre. Pour mieux faire juger à l'œil la diminution de l'armée, j'ai supposé que les corps du Prince Jérôme et du Maréchal Davoust qui avaient été détachés sur Minsk et Mabilow et qui rejoignent vers Orscha et Witebsk, avaient toujours marché avec l'armée.



Auget par Regnier, 8, Par. 5^{de} Marie St 6^{de} à Paris.

Imp. Lith. Regnier et Dourdet.

Charles Minard's Flow Map of Napoleon's Russian Campaign of 1812



Common Data Visualizations

- Summary statistics
- Box plots
- Histograms
- Scatter plots


Big Data Visualizations

- Large n
 - Computationally expensive to render
 - Dense/complex
 - Address via subsampling or parallelization
- Large d
 - Difficult to represent more than a few dimensions
 - Address via dimensionality reduction = learning a latent (typically lower-dimensional) representation


Feature Elimination



Dimensionality Reduction




Option A




Option B

Which projection do you prefer?

Goal: minimize the reconstruction error





Option A



Option B

Goal: maximize the variance of the projections



Centering the Data

- To be consistent, we will constrain principal components to be orthonormal vectors (orthogonal unit vectors) that begin at the origin
- Preprocess data to be centered around the origin:

$$\mu = \frac{1}{n} \sum_{i=1}^n x^{(i)}$$

$$\tilde{x}^{(i)} = x^{(i)} - \mu$$

$$X = \begin{bmatrix} \tilde{x}^{(1)T} \\ \vdots \\ \tilde{x}^{(n)T} \end{bmatrix} \in \mathbb{R}^{n \times d}$$

Reconstruction Error

- The projection of $\tilde{x}^{(i)}$ onto a vector v is

$$z^{(i)} = \left(\frac{v^T \tilde{x}^{(i)}}{\|v\|_2} \right) \frac{v}{\|v\|_2}$$

Length of projection

Direction of projection

Reconstruction Error

- The projection of $\tilde{x}^{(i)}$ onto a unit vector v is

$$z^{(i)} = (v^T \tilde{x}^{(i)})v$$

$$\begin{aligned}\hat{v} = \operatorname{argmin}_{v: \|v\|_2^2=1}(\text{the resstruction error}) &= \operatorname{argmin}_{v: \|v\|_2^2=1} \sum_{i=1}^n \|\tilde{x}^{(i)} - z^{(i)}\|_2^2 \\ &= \operatorname{argmin}_{v: \|v\|_2^2=1} \sum_{i=1}^n \|\tilde{x}^{(i)} - (v^T \tilde{x}^{(i)})v\|_2^2 \\ &= \operatorname{argmin}_{v: \|v\|_2^2=1} \sum_{i=1}^n \|\tilde{x}^{(i)}\|_2^2 - (v^T \tilde{x}^{(i)})^2\end{aligned}$$

Minimizing the Reconstruction Error

$$\begin{aligned}\hat{v} &= \operatorname{argmin}_{v: \|v\|_2^2=1} \sum_{i=1}^n \|\tilde{x}^{(i)}\|_2^2 - (v^T \tilde{x}^{(i)})^2 = \operatorname{argmin}_{v: \|v\|_2^2=1} - \sum_{i=1}^n (v^T \tilde{x}^{(i)})^2 = \operatorname{argmax}_{v: \|v\|_2^2=1} \sum_{i=1}^n (v^T \tilde{x}^{(i)})^2 \\ &= \operatorname{argmax}_{v: \|v\|_2^2=1} \sum_{i=1}^n v^T \tilde{x}^{(i)} \tilde{x}^{(i)T} v = \operatorname{argmax}_{v: \|v\|_2^2=1} v^T \left(\sum_{i=1}^n \tilde{x}^{(i)} \tilde{x}^{(i)T} \right) v \\ &= \operatorname{argmax}_{v: \|v\|_2^2=1} v^T (X^T X) v = \operatorname{argmax}_{v: \|v\|_2^2=1} v^T C_X v\end{aligned}$$

C_X : covariance matrix

Maximizing the Variance

$$\hat{v} = \operatorname{argmax}_{v: \|v\|_2^2=1} v^T C_X v$$

$$L(v, \lambda) = v^T C_X v + \lambda(\|v\|_2^2 - 1) = v^T C_X v + \lambda(v^T v - 1)$$

$$\frac{\partial L}{\partial v} = C_X v - \lambda v = 0$$

$\Rightarrow \lambda$ is eigenvalue of C_X and its corresponding eigenvector is v

Maximizing the Variance

$$\hat{v} = \operatorname{argmax}_{v: \|v\|_2^2=1} v^T C_X v$$

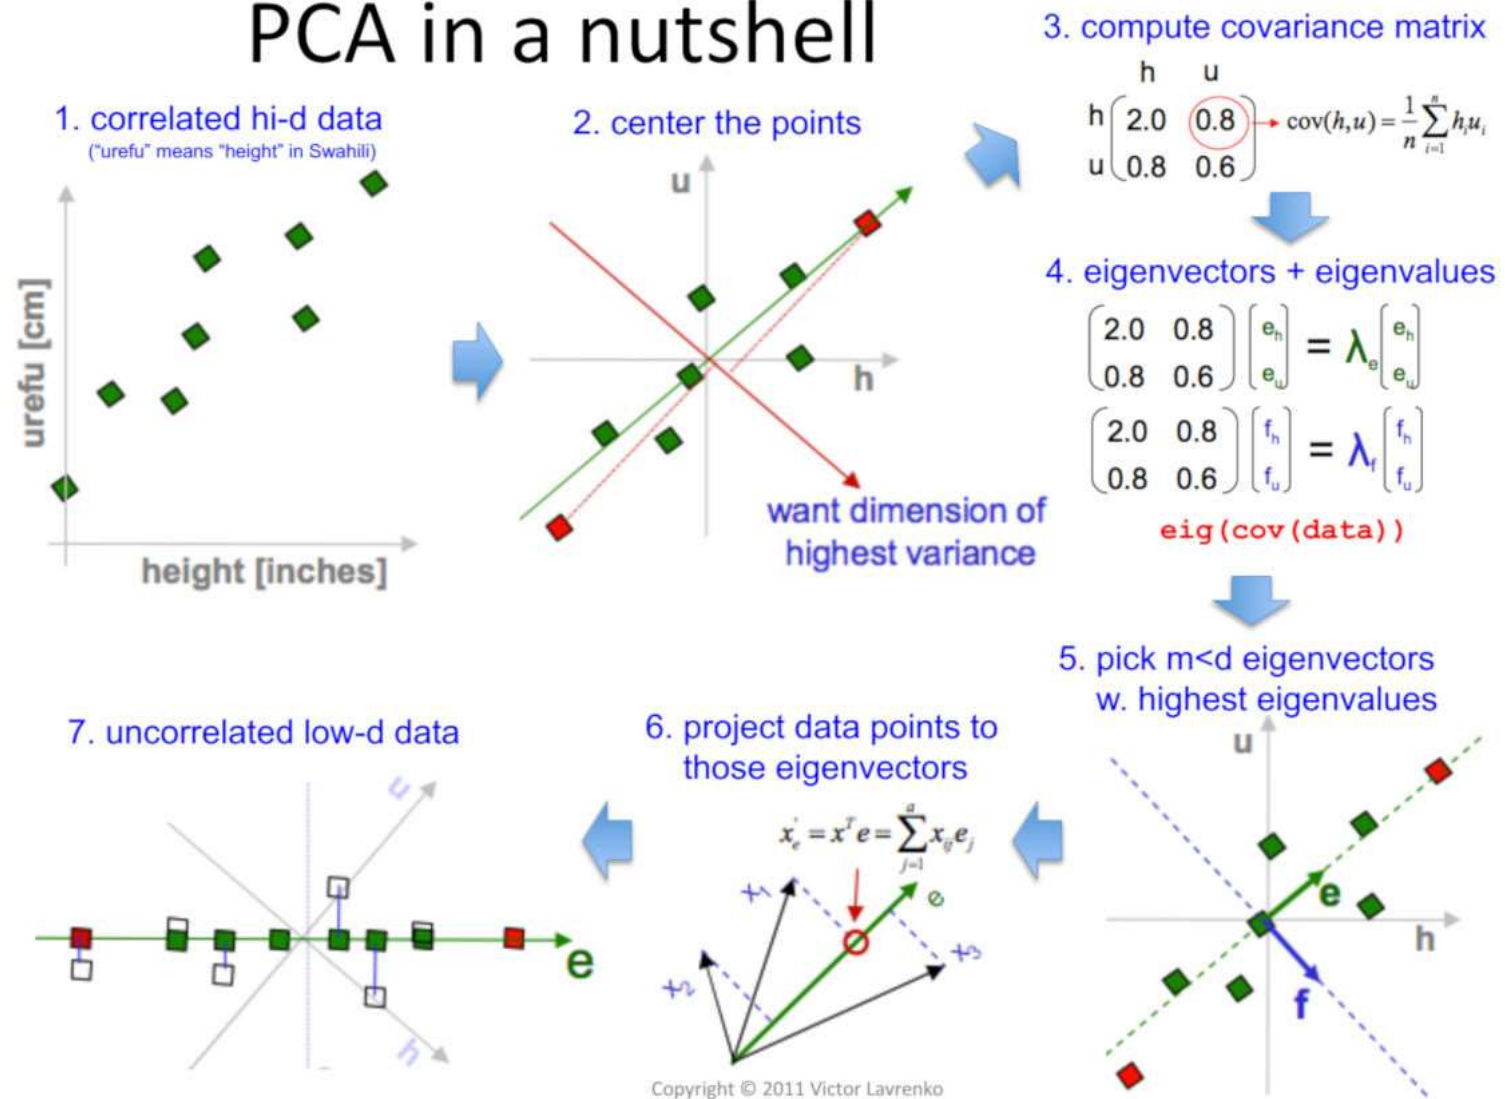
$$C_X v = \lambda v \Rightarrow v^T C_X v = \lambda v^T v = \lambda$$

- The first principal component is the eigenvector \hat{v}_1 that corresponds to the largest eigenvalue λ_1
- The second principal component is the eigenvector \hat{v}_2 that corresponds to the second largest eigenvalue λ_2
- ...
- λ_i is a measure of how much variance falls along \hat{v}_i

PCA Algorithm

- Input: $D = \{(x^{(i)}, y^{(i)})\}_{i=1}^n, k$
 1. Center the data
 - Optionally, normalize the data by features so that all features are of the same scale
 2. Compute the covariance matrix $C_X = X^T X$
 3. Collect the top k eigenvectors (corresponding to the largest eigenvalues), $P \in \mathbb{R}^{n \times k}$
 4. Project the data into the space defined by P , $Z = XP$
- Output: Z , the latent representation (PCA scores)

PCA in a nutshell




Choosing the number of PCs

- Define a percentage of explained variance for the i th PC:


$$\frac{\lambda_i}{\sum_{j=1}^n \lambda_j}$$

- Select all PCs above some threshold of explained variance, e.g., 5%
- Keep selecting PCs until the total explained variance exceeds some threshold, e.g., 90%
- Evaluate on some downstream metric

PCA Example: MNIST Digits




PCA Example: MNIST Digits



Shortcomings of PCA

- Sometime we don't care about variance
- PCA only find linear combination of our features
- Principal components are expensive to compute
- Interpretability
- Principal components are orthonormal



PCA Algorithm: Computational Cost

- Input: $D = \{(x^{(i)}, y^{(i)})\}_{i=1}^n, k$
 1. Center the data
 - Optionally, normalize the data by features so that all features are of the same scale
 2. Compute the covariance matrix $C_X = X^T X \rightarrow O(nd^2)$
 3. Collect the top k eigenvectors (corresponding to the largest eigenvalues), $P \in \mathbb{R}^{n \times k} \rightarrow O(d^3)$
 4. Project the data into the space defined by P , $Z = XP \rightarrow O(ndk)$
- Output: Z , the latent representation (PCA scores)

Nonlinear Dimensionality Reduction

- Autoencoders (1987)
- Kernel PCA (1999)
- Locally linear embedding (2000)
- Isomap (2000)
- Laplacian Eigenmaps (2003)
- t-SNE (2008)
- ... many others

PCA: Large n , Small d

- Assume $O(d^3)$ computation and $O(d^2)$ storage is possible on a single machine
 - We **can** store and compute the eigenvalues of $X^T X$
 - We **cannot** compute $X^T X$
 - We **cannot** store X
- Approach: basically the same as distributed linear regression
 1. Center the data in a distributed way
 2. Store the rows of X across different machines
 3. Compute $X^T X$ as the sum of outer products

Distributed Centering of the Data

Worker	$\begin{bmatrix} \leftarrow & x^{(1)T} & \rightarrow \\ \leftarrow & x^{(4)T} & \rightarrow \\ \vdots & \vdots & \vdots \end{bmatrix}$	$\begin{bmatrix} \leftarrow & x^{(2)T} & \rightarrow \\ \leftarrow & x^{(3)T} & \rightarrow \\ \vdots & \vdots & \vdots \end{bmatrix}$	$\begin{bmatrix} \leftarrow & x^{(5)T} & \rightarrow \\ \leftarrow & x^{(7)T} & \rightarrow \\ \vdots & \vdots & \vdots \end{bmatrix}$
Reduce	$\mu = \frac{1}{n} \sum_{i=1}^n x^{(i)}$		
Map	$\tilde{x}^{(i)} = x^{(i)} - \mu$	$\tilde{x}^{(i)} = x^{(i)} - \mu$	$\tilde{x}^{(i)} = x^{(i)} - \mu$

$O(nd)$ distributed
storage (total)

$O(d)$ local work $O(d)$ local storage

$O(nd)$ distributed
work (total) $O(nd)$ distributed
storage (total)

$O(d)$ communication

Distributed Eigendecomposition of $X^T X$

$$X^T X = \begin{bmatrix} \uparrow & \uparrow & \uparrow \\ x^{(1)} & \cdots & x^{(n)} \\ \downarrow & \downarrow & \downarrow \end{bmatrix} \begin{bmatrix} \leftarrow & x^{(1)} & \rightarrow \\ \vdots & \vdots & \vdots \\ \leftarrow & x^{(n)} & \rightarrow \end{bmatrix} = \sum_{i=1}^n x^{(i)} x^{(i)T}$$

Distributed Eigendecomposition of $X^T X$

Worker	$\begin{bmatrix} \leftarrow & x^{(1)T} & \rightarrow \\ \leftarrow & x^{(4)T} & \rightarrow \\ \vdots & \vdots & \vdots \end{bmatrix}$	$\begin{bmatrix} \leftarrow & x^{(2)T} & \rightarrow \\ \leftarrow & x^{(3)T} & \rightarrow \\ \vdots & \vdots & \vdots \end{bmatrix}$	$\begin{bmatrix} \leftarrow & x^{(5)T} & \rightarrow \\ \leftarrow & x^{(7)T} & \rightarrow \\ \vdots & \vdots & \vdots \end{bmatrix}$	$O(nd)$ distributed storage (total)
Map	$\tilde{x}^{(i)} \tilde{x}^{(i)T}$	$\tilde{x}^{(i)} \tilde{x}^{(i)T}$	$\tilde{x}^{(i)} \tilde{x}^{(i)T}$	$O(nd^2)$ distributed work (total) $O(d^2)$ local storage
Reduce	$eigh\left(\sum_{i=1}^n \tilde{x}^{(i)} \tilde{x}^{(i)T}\right)$			$O(d^3)$ local work $O(d^2)$ local storage

$O(dk)$ communication

Distributed Computation of PCA Scores

Worker	$\begin{bmatrix} \leftarrow & x^{(1)T} & \rightarrow \\ \leftarrow & x^{(4)T} & \rightarrow \\ \vdots & \vdots & \vdots \end{bmatrix}$	$\begin{bmatrix} \leftarrow & x^{(2)T} & \rightarrow \\ \leftarrow & x^{(3)T} & \rightarrow \\ \vdots & \vdots & \vdots \end{bmatrix}$	$\begin{bmatrix} \leftarrow & x^{(5)T} & \rightarrow \\ \leftarrow & x^{(7)T} & \rightarrow \\ \vdots & \vdots & \vdots \end{bmatrix}$	$O(nd)$ distributed storage (total)
Map	$P\tilde{x}^{(i)}$	$P\tilde{x}^{(i)}$	$P\tilde{x}^{(i)}$	$O(ndk)$ $O(nk)$ distributed local work (total) storage

PCA: Large n , Large k

- Now, $O(d^3)$ computation and $O(d^2)$ storage is not possible on a single machine
 - We cannot store and compute the eigenvalues of $X^T X$
 - We cannot compute $X^T X$
 - We cannot store X
- Idea: use a different algorithm!
 - Turn to an iterative method for computing eigenvectors

PCA: Large n , Large k

- Now, $O(d^3)$ computation and $O(d^2)$ storage is not possible on a single machine
 - We cannot store and compute the eigenvalues of $X^T X$
 - We cannot compute $X^T X$
 - We cannot store X
- Idea: use a different algorithm!
 - Turn to an iterative method for computing ~~eigenvalues~~ the eigenvector associated with the largest eigenvalue ($k = 1$) \rightarrow power iteration

Power Iteration

- Fact: $A = X^T X$ is “diagonalizable”, i.e., any d-dimensional vector can be written as a linear combination of A ’s eigenvectors:

$$b = c_1 v_1 + c_2 v_2 + \cdots + c_d v_d$$

- This follows because $X^T X$ is real and symmetric
- Assume $A = X^T X$ has one eigenvalue that is strictly larger than the others:

$$\lambda_1 \geq \lambda_2 \geq \cdots \geq \lambda_d$$

Power Iteration

- Input: $A = X^T X$
- Initialize $b^{(0)}$ randomly and set $t = 0$
- While **not converged**
 - Update the vector b :

$$b^{(t+1)} = \frac{Ab^{(t)}}{\|Ab^{(t)}\|_2}$$

- Increment t : $t = t + 1$
- Output: $b^{(t)}$, the eigenvector corresponding to the largest eigenvalue of A

Power Iteration

$$b^{(0)} = c_1 v_1 + c_2 v_2 + \cdots + c_d v_d$$

$$Ab^{(0)} = c_1 A v_1 + c_2 A v_2 + \cdots + c_d A v_d$$

$$= c_1 \lambda_1 v_1 + c_2 \lambda_2 v_2 + \cdots + c_d \lambda_d v_d$$

$$A(Ab^{(0)}) = c_1 \lambda_1 A v_1 + c_2 \lambda_2 A v_2 + \cdots + c_d \lambda_d A v_d$$

$$= c_1 \lambda_1^2 v_1 + c_2 \lambda_2^2 v_2 + \cdots + c_d \lambda_d^2 v_d$$

$$A^t b^{(0)} = c_1 \lambda_1^t v_1 + c_2 \lambda_2^t v_2 + \cdots + c_d \lambda_d^t v_d$$

$$= \lambda_1^t \left(c_1 v_1 + c_2 \left(\frac{\lambda_2}{\lambda_1} \right)^t v_2 + \cdots + c_d \left(\frac{\lambda_d}{\lambda_1} \right)^t v_d \right) \xrightarrow{t \rightarrow \infty} \lambda_1^t c_1 v_1$$

Power Iteration

- Input: $A = X^T X$
- Initialize $b^{(0)}$ randomly and set $t = 0$
- While **not converged**
 - Update the vector b :

$$b^{(t+1)} = \frac{X^T X b^{(t)}}{\|X^T X b^{(t)}\|_2}$$

- Increment t : $t = t + 1$
- Output: $b^{(t)}$, the eigenvector corresponding to the largest eigenvalue of A

Power Iteration

- Input: $A = X^T X$
- Initialize $b^{(0)}$ randomly and set $t = 0$
- While **not converged**

- Update the vector b :

$$b^{(t+1)} = X^T X b^{(t)} = \left(\sum_{i=1}^n x^{(i)} x^{(i)T} \right) b^{(t)}$$
$$b^{(t+1)} = \frac{b^{(t+1)}}{\|b^{(t+1)}\|_2}$$

- Increment t : $t = t + 1$
- Output: $b^{(t)}$, the eigenvector corresponding to the largest eigenvalue of A

Power Iteration

- Input: $A = X^T X$
- Initialize $b^{(0)}$ randomly and set $t = 0$
- While **not converged**

- Update the vector b :

$$b^{(t+1)} = X^T X b^{(t)} = \sum_{i=1}^n x^{(i)} \left(x^{(i)T} b^{(t)} \right)$$
$$b^{(t+1)} = \frac{b^{(t+1)}}{\|b^{(t+1)}\|_2}$$

- Increment t : $t = t + 1$
- Output: $b^{(t)}$, the eigenvector corresponding to the largest eigenvalue of A

Power Iteration

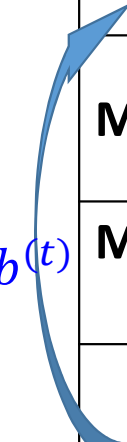
- Input: $A = X^T X$
- Initialize $b^{(0)}$ randomly and set $t = 0$
- While **not converged**

- Update the vector b :

$$b^{(t+1)} = X^T X b^{(t)} = \sum_{i=1}^n x^{(i)} \left(\beta_i^{(t)} \right)$$
$$b^{(t+1)} = \frac{b^{(t+1)}}{\|b^{(t+1)}\|_2}$$

- Increment t : $t = t + 1$
- Output: $b^{(t)}$, the eigenvector corresponding to the largest eigenvalue of A

Distributed Power Iteration



Worker	$\begin{bmatrix} \leftarrow & x^{(1)T} & \rightarrow \\ \leftarrow & x^{(4)T} & \rightarrow \\ \vdots & \vdots & \vdots \end{bmatrix}$	$\begin{bmatrix} \leftarrow & x^{(2)T} & \rightarrow \\ \leftarrow & x^{(3)T} & \rightarrow \\ \vdots & \vdots & \vdots \end{bmatrix}$	$\begin{bmatrix} \leftarrow & x^{(5)T} & \rightarrow \\ \leftarrow & x^{(7)T} & \rightarrow \\ \vdots & \vdots & \vdots \end{bmatrix}$	$O(nd)$ distributed storage (total)
Map	$\beta_i^{(t)} = \tilde{x}^{(i)T} b^{(t)}$	$\beta_i^{(t)} = \tilde{x}^{(i)T} b^{(t)}$	$\beta_i^{(t)} = \tilde{x}^{(i)T} b^{(t)}$	$O(nd)$ distributed work (total) $O(n)$ distributed storage
Map	$\tilde{x}^{(i)} \beta_i^{(t)}$	$\tilde{x}^{(i)} \beta_i^{(t)}$	$\tilde{x}^{(i)} \beta_i^{(t)}$	$O(nd)$ distributed work (total) $O(d)$ local storage
Reduce	$b^{(t+1)} = \sum_{i=1}^n \tilde{x}^{(i)} \beta_i^{(t)} / \left\ \sum_{i=1}^n \tilde{x}^{(i)} \beta_i^{(t)} \right\ _2$			$O(d)$ local work $O(d)$ local storage

$O(d)$ communication To find more than one eigenvector, we need to use similar (but slightly more complicated methods), e.g., PySpark's MLlib uses Krylov subspace methods 41

Summary

- Distributed PCA very similar to distributed linear regression
 - If d is small, simply distribute the storage and computation of $X^T X$
 - If d is large, use iterative methods (e.g., power iteration) to compute eigenvector-eigenvalue pairs